
Distributed Transaction Management

Transaction Concepts and Models

Contents

- Definition of a Transaction
 - Termination Conditions
 - Characterization and Formalization
- Properties of Transactions
 - ACID: Atomicity, Consistency, Isolation, Durability
- Architecture Revisited

Transaction

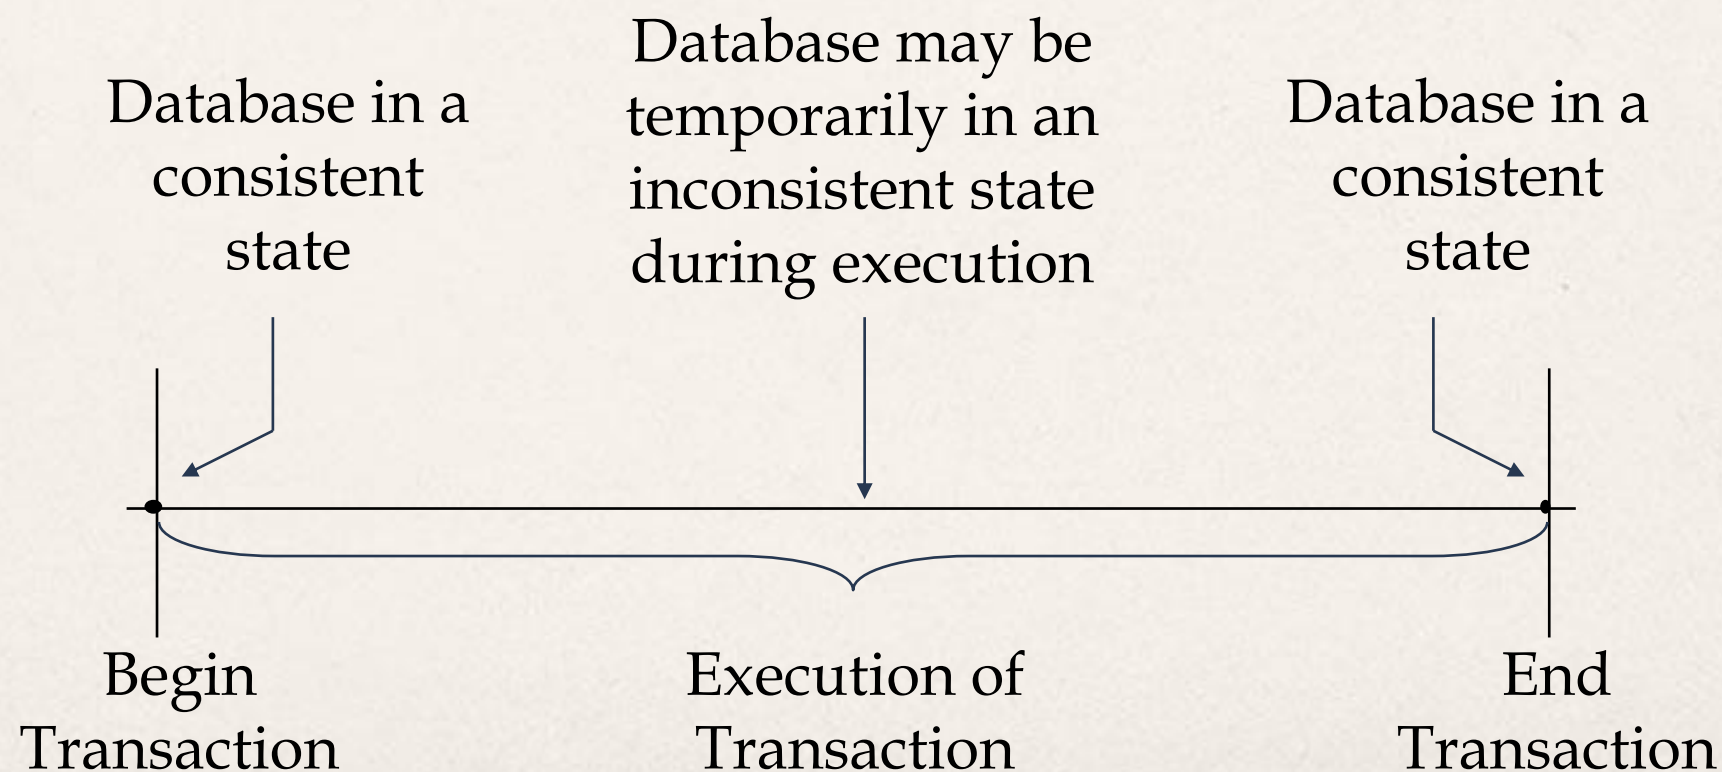
A transaction is a collection of actions that make consistent transformations of system states while **preserving system consistency**.

→ **concurrency transparency**

concurrency control

→ **failure transparency**

reliability



Transaction Example

A Simple SQL Query

Begin_transaction Salary_Update

begin

EXEC SQL	UPDATE	PAY
	SET	SAL = SAL*1.1

end

Example Database

Consider an airline reservation example with the relations:

FLIGHT(FNO, DATE, SRC, DEST, STSOLD, CAP)

CUST(CNAME, ADDR, BAL)

FC(FNO, DATE, CNAME, SPECIAL)

Transaction Example

SQL Notation

Begin_transaction Reservation

begin

input(flight_no, date, customer_name);

EXEC SQL UPDATE FLIGHT

SET STSOLD = STSOLD + 1

WHERE FNO = flight_no AND DATE = date;

EXEC SQL INSERT

INTO FC(FNO, DATE, CNAME, SPECIAL);

VALUES (flight_no, date, customer_name, **null**);

output("reservation completed")

end

Termination Conditions

Begin_transaction Reservation

begin

input(flight_no, date, customer_name);

 EXEC SQL SELECT STSOLD,CAP

 INTO temp1,temp2

 FROM FLIGHT

 WHERE FNO = flight_no AND DATE = date;

if temp1 = temp2 **then**

output("no free seats");

Abort

else

 EXEC SQL UPDATE FLIGHT

 SET STSOLD = STSOLD + 1

 WHERE FNO = flight_no AND DATE = date;

 EXEC SQL INSERT

 INTO FC(FNO, DATE, CNAME, SPECIAL);

 VALUES (flight_no, date, customer_name, **null**);

Commit

output("reservation completed")

endif

end

Characterization

- Read set (RS)
 - The set of data items that are read by a transaction
- Write set (WS)
 - The set of data items whose values are changed by this transaction
- Base set (BS)
 - $RS \cup WS$
- $RS[\text{Reservation}] = \{\text{FLIGHT.STSOLD}, \text{FLIGHT.CAP}, \text{FLIGHT.FNO}, \text{FLIGHT.DATE}\}$
- $WS[\text{Reservation}] = \{\text{FLIGHT.STSOLD}, \text{FC.FNO}, \text{FC.DATE}, \text{FC.CNAME}, \text{FC.SPECIAL}\}$

Formalization

Let

$\Rightarrow O_{ij}(x)$ be some operation O_j of transaction T_i operating on entity x , where $O_j \in \{\text{read}, \text{write}\}$ and O_j is atomic

$\Rightarrow OS_i = \cup_j O_{ij}$

$\Rightarrow N_i \in \{\text{abort}, \text{commit}\}$

Transaction T_i is a partial order $T_i = \{\Sigma_i, <_i\}$ where

- ① $\Sigma_i = OS_i \cup \{N_i\}$
- ② For any two operations $O_{ij}, O_{ik} \in OS_i$, if $O_{ij} = R(x)$ and $O_{ik} = W(x)$ for any data item x , then either $O_{ij} <_i O_{ik}$ or $O_{ik} <_i O_{ij}$
- ③ For all $O_{ij} \in OS_i$, $O_{ij} <_i N_i$

Example

Consider a transaction T :

Read(x)

Read(y)

$x \leftarrow x + y$

Write(x)

Commit

Then


$\Sigma = \{R(x), R(y), W(x), C\}$

$< = \{(R(x), W(x)), (R(y), W(x)), (W(x), C), (R(x), C), (R(y), C)\}$

DAG Representation

Assume

$$< = \{(R(x), W(x)), (R(y), W(x)), (W(x), C), (R(x), C), (R(y), C)\}$$



Properties of Transactions

A TOMICITY

→ all or nothing

C ONSISTENCY

→ no violation of integrity constraints

I SOLATION

→ concurrent changes invisible

D URABILITY

→ committed updates persist

Atomicity

- Either **all or none** of the transaction's operations are performed.
- Atomicity requires that if a transaction is interrupted by a failure, its partial results must be **undone**.
- The activity of preserving the transaction's atomicity in presence of transaction aborts due to input errors, system overloads, or deadlocks is called **transaction recovery**.
- The activity of ensuring atomicity in the presence of system crashes is called **crash recovery**.

Consistency

- Internal consistency
 - A transaction which executes **alone** against a **consistent** database leaves it in a consistent state.
 - Transactions do not violate database integrity constraints.
- Transactions are **correct** programs

Isolation

- Concurrent changes invisible (incomplete results)
 - An executing transaction cannot reveal its results to other concurrent transactions before its commitment.
 - Necessary to avoid cascading aborts.
- Serializability
 - If several transactions are executed concurrently, the results must be the same as if they were executed serially in some order (**correctness criterion**).
- Concurrent changes invisible \Rightarrow serializability

Isolation Example

- Consider the following two transactions:

T_1 : Read(x)
 $x \leftarrow x+1$
 Write(x)
 Commit

T_2 : Read(x)
 $x \leftarrow x+1$
 Write(x)
 Commit

- Possible execution sequences:

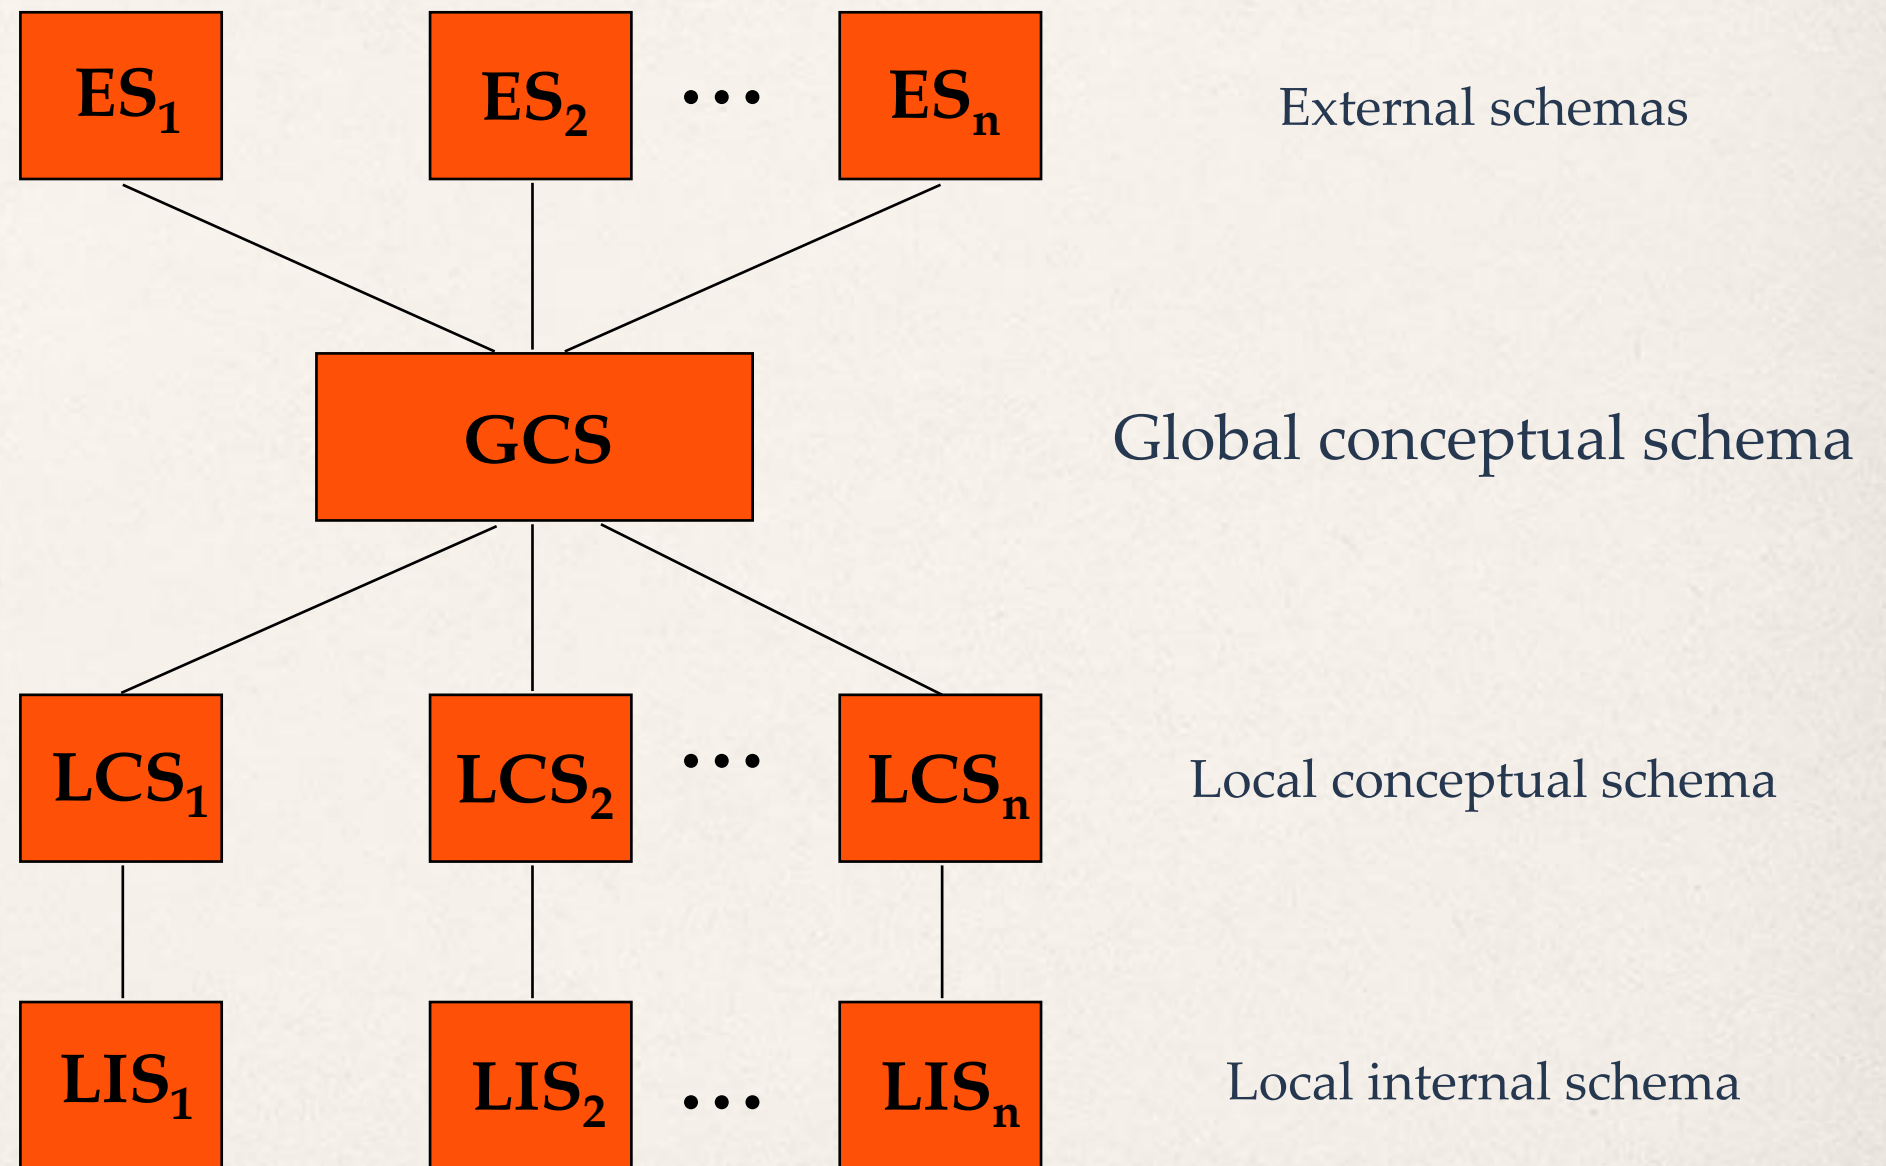
T_1 : Read(x)
 T_1 : $x \leftarrow x+1$
 T_1 : Write(x)
 T_1 : Commit
 T_2 : Read(x)
 T_2 : $x \leftarrow x+1$
 T_2 : Write(x)
 T_2 : Commit

T_1 : Read(x)
 T_1 : $x \leftarrow x+1$
 T_2 : Read(x)
 T_1 : Write(x) ×
 T_2 : $x \leftarrow x+1$
 T_2 : Write(x)
 T_1 : Commit
 T_2 : Commit

Durability

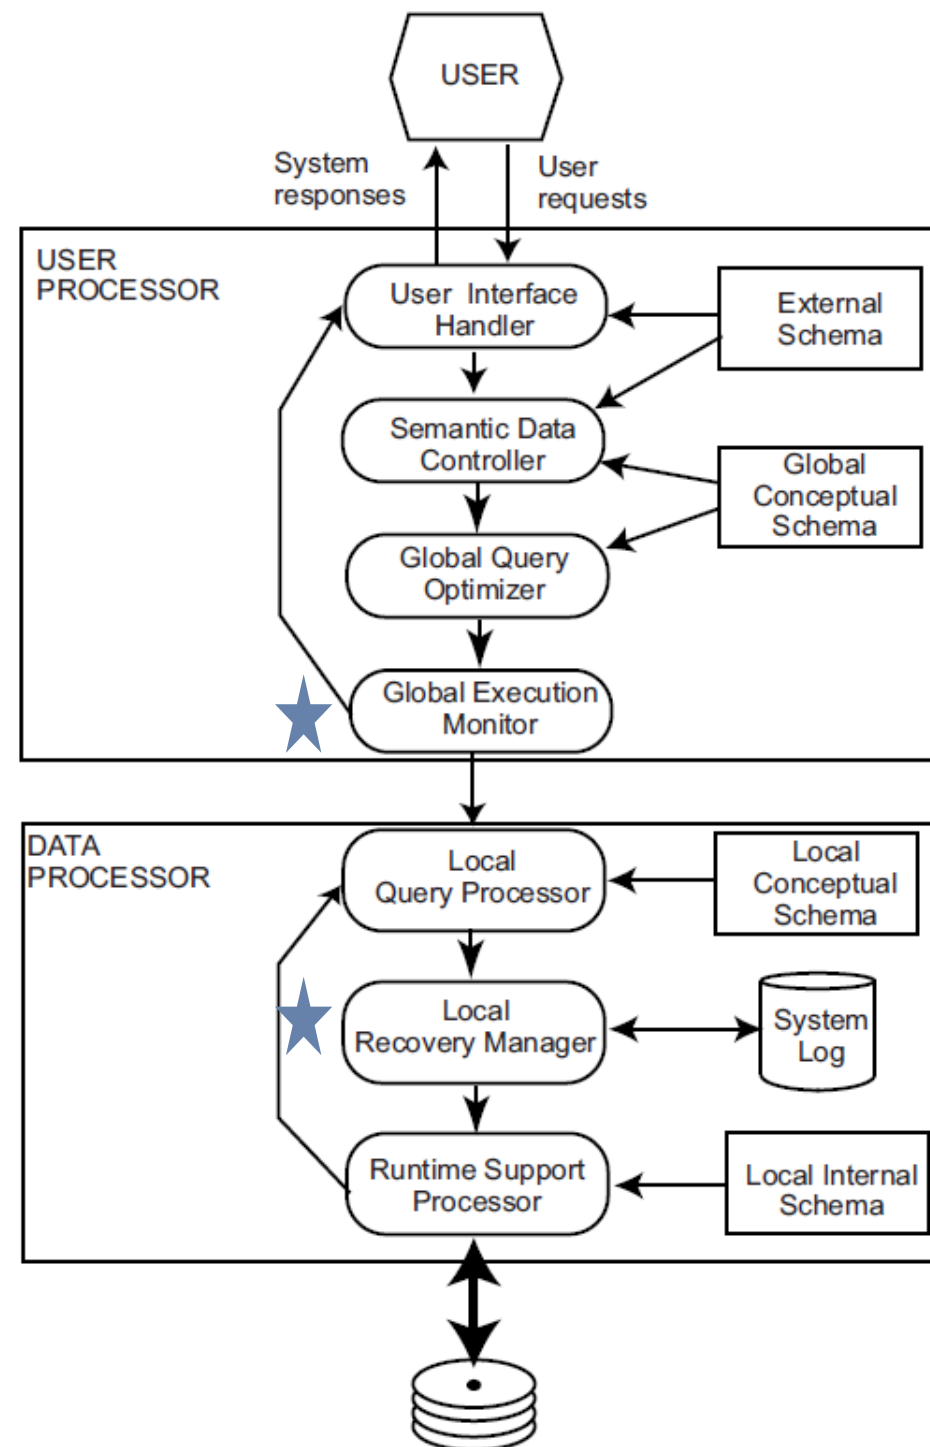
- Once a transaction commits, the system must guarantee that the results of its operations will never be lost, in spite of subsequent failures.
- Database recovery

Distributed Database Reference Architecture

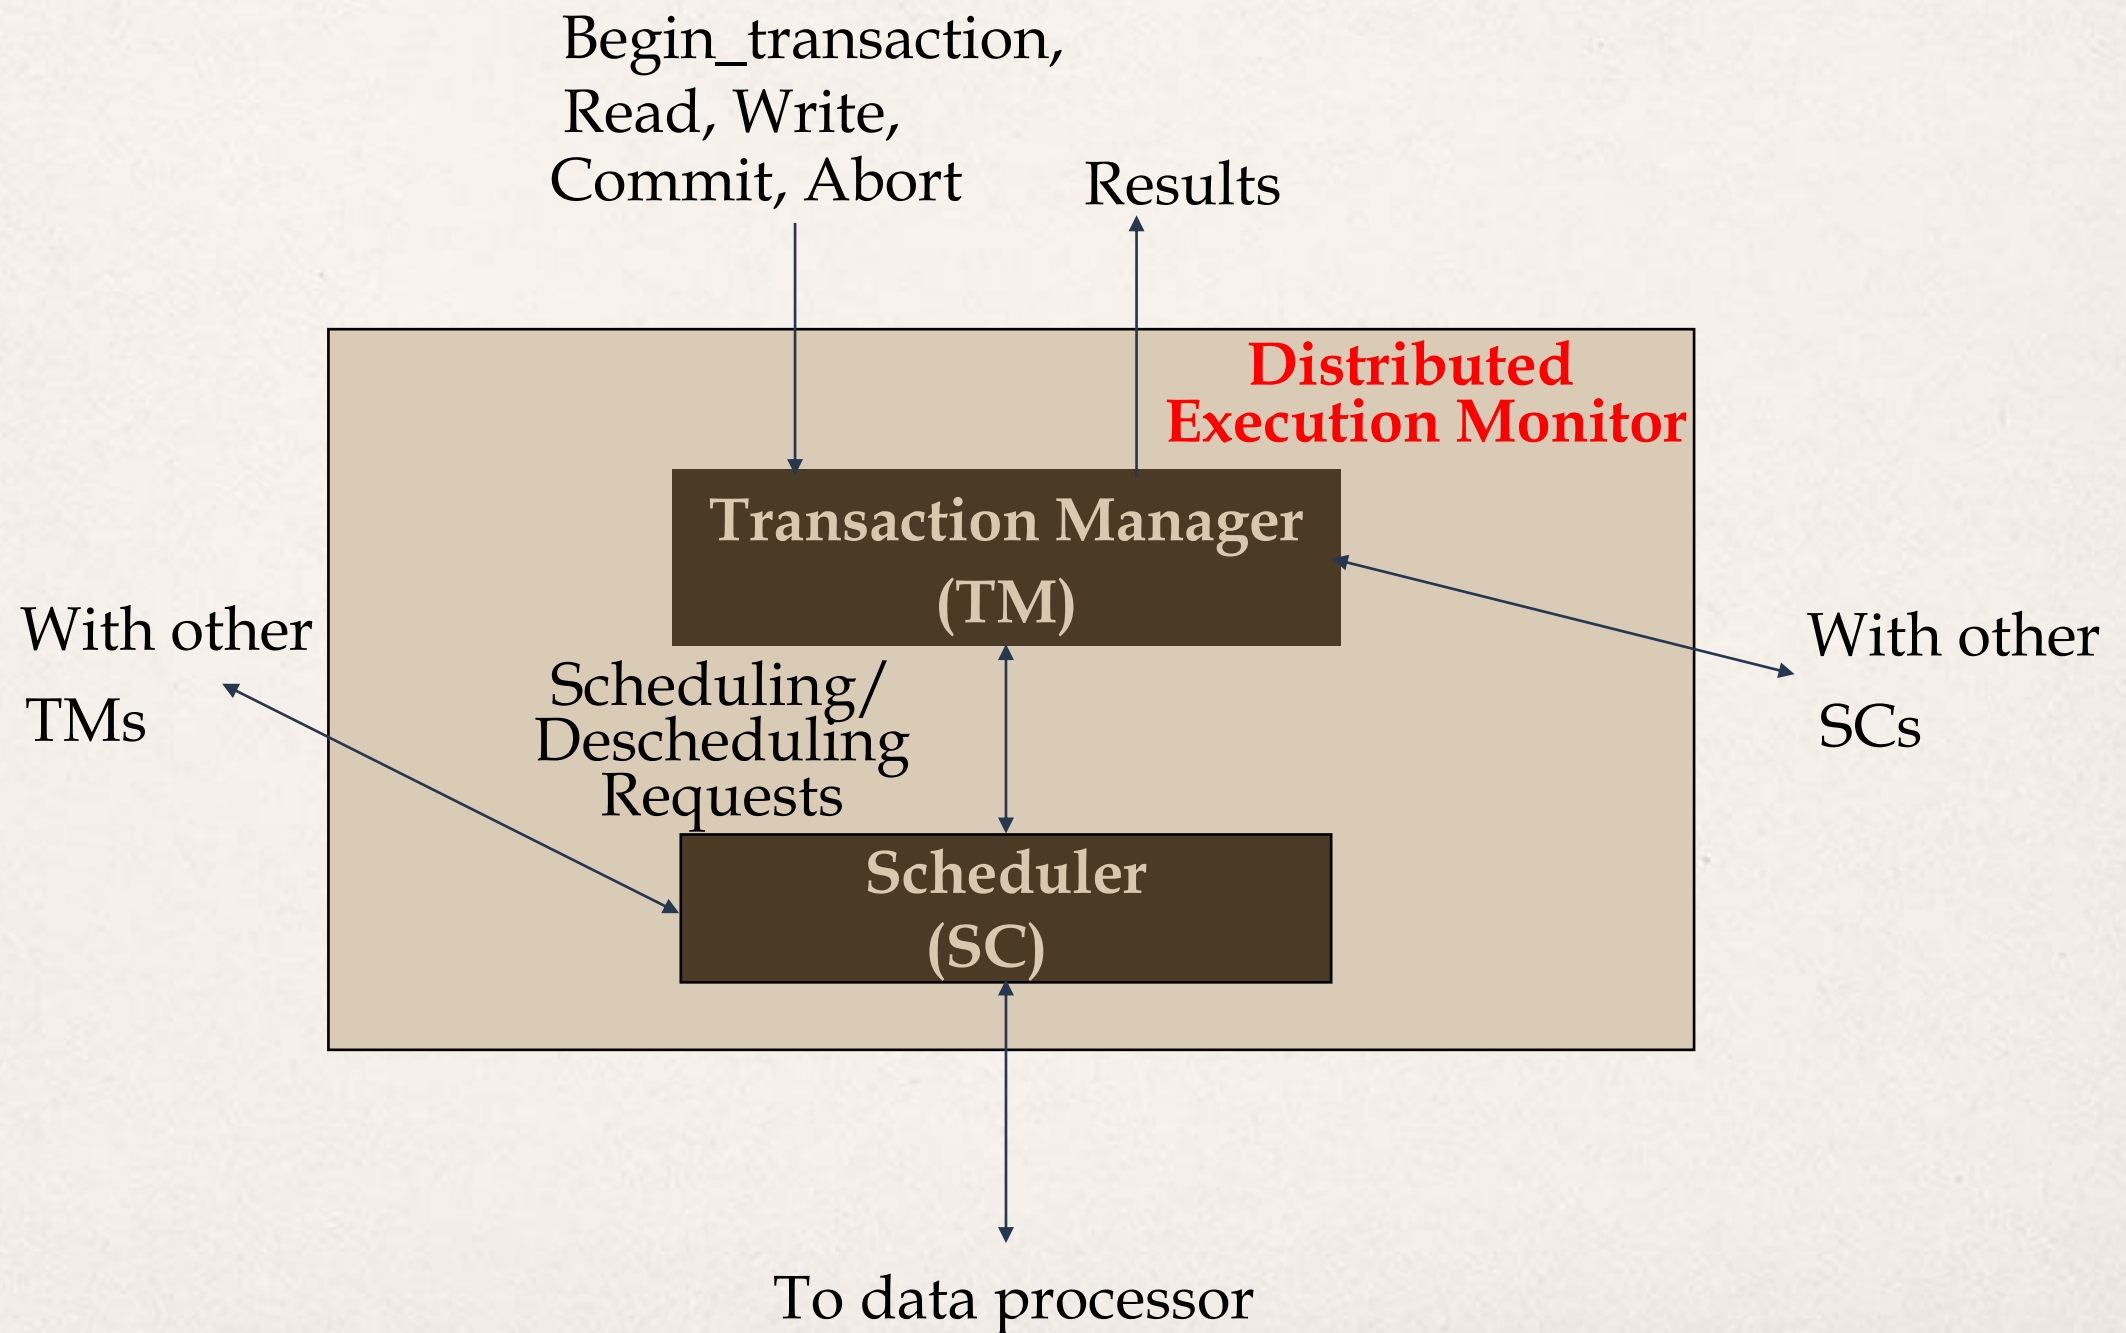


Data organizational view

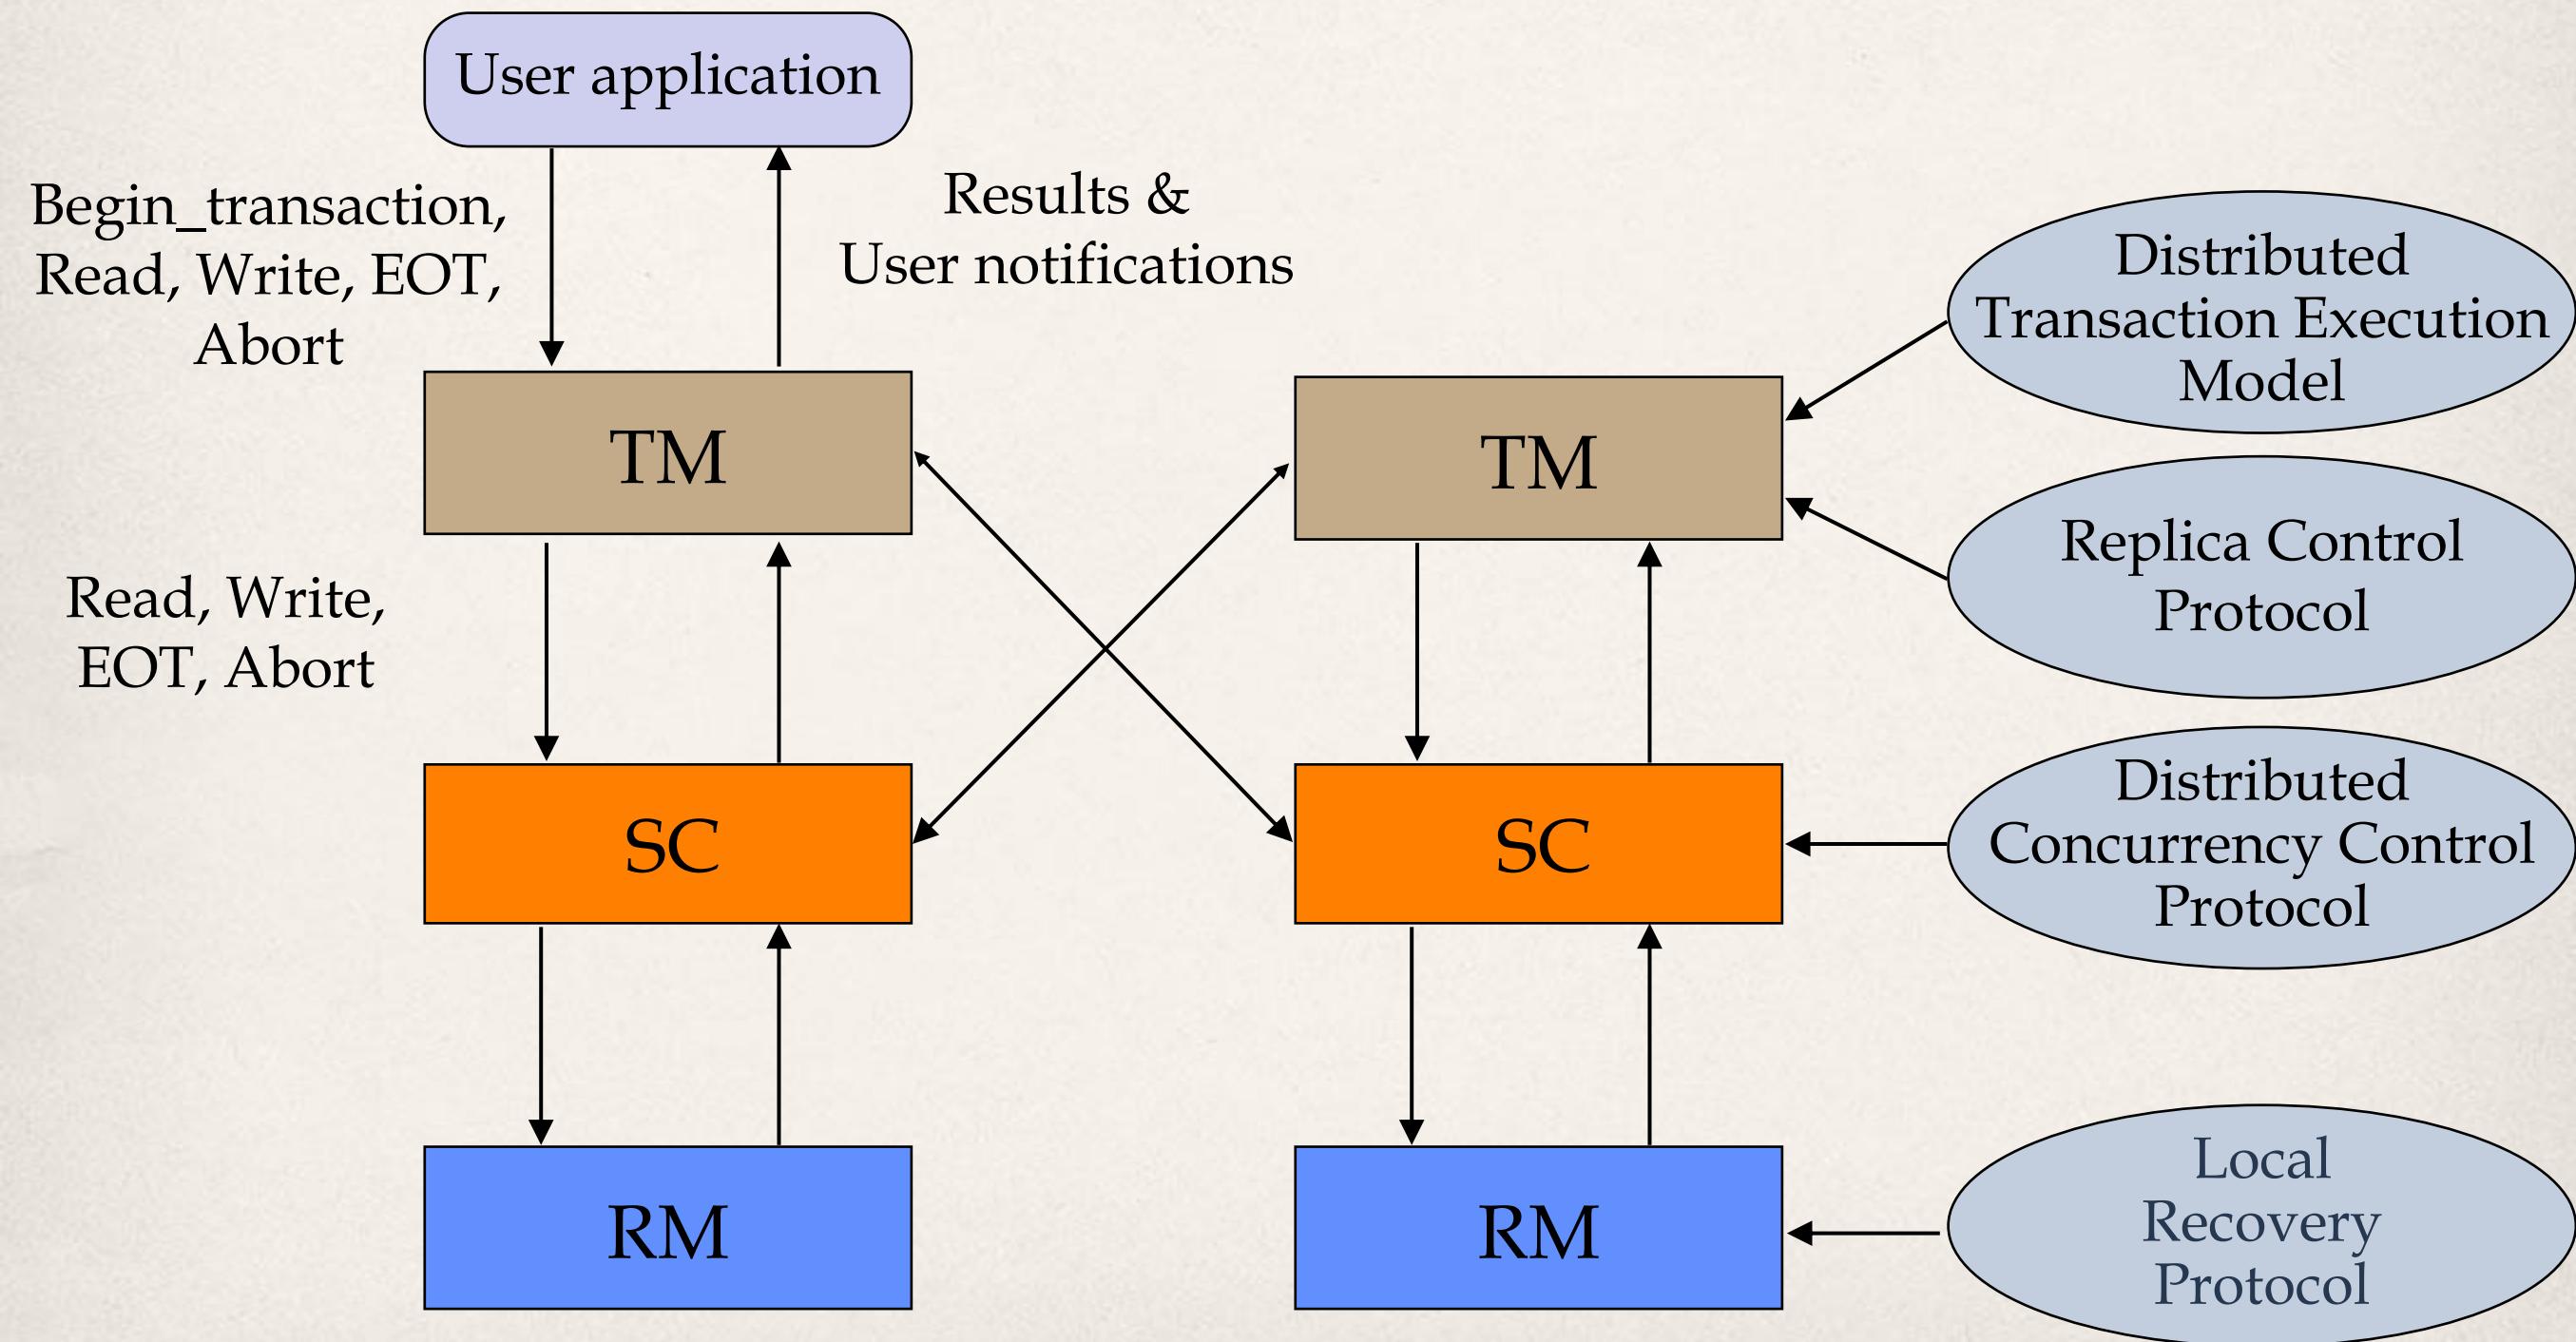
Components of a Distributed DBMS



Architecture Revisited



Distributed Transaction Execution



Distributed Transaction Management

Transaction Concepts and Models

Consistency Degrees

- Degree 3

- T does not overwrite dirty data of other transactions
- T does not commit any writes before EOT
- T does not read dirty data from other transactions
- Other transactions do not dirty any data read by T before T completes.

- ◆ Dirty data refers to data values that have been updated by a transaction prior to its commitment

Consistency Degrees (cont'd)

- Degree 2

- T does not overwrite dirty data of other transactions
- T does not commit any writes before EOT
- T does not read dirty data from other transactions

- Degree 1

- T does not overwrite dirty data of other transactions
- T does not commit any writes before EOT

- Degree 0

- Transaction T does not overwrite dirty data of other transactions

SQL-92 Isolation Levels

- Defined on the basis of what ANSI call phenomena:

- ☞ Dirty read

- ◆ T_1 modifies x which is then read by T_2 before T_1 terminates. If T_1 aborts, T_2 has read a value which never exists in the database.

- ☞ Non-repeatable (fuzzy) read

- ◆ T_1 reads x ; T_2 then modifies or deletes x and commits. If T_1 tries to read x again, it reads a different value or it can't find it.

- ☞ Phantom

- ◆ T_1 searches the database according to a predicate while T_2 inserts new tuples that satisfy the predicate.

SQL-92 Isolation Levels (cont'd)

- Anomaly Serializable
 - ☞ None of the phenomena are possible.
- Repeatable Read
 - ☞ Only phantoms possible.
- Read Committed
 - ☞ Fuzzy reads and phantoms are possible, but dirty reads are not.
- Read Uncommitted
 - ☞ For transactions operating at this level, all three phenomena are possible.

Types of Transactions

- Based on

- Timing

- ◆ On-line (short-life) vs batch (long-life)

- Organization of read and write actions

- ◆ General
 - ◆ Two-step
 - ◆ Restricted
 - ◆ Action model

- **Structure**

- ◆ Flat (or simple) transactions
 - ◆ Nested transactions

Transaction Structure

- Flat transaction

→ Consists of a sequence of **primitive** operations embraced between **begin** and **end** markers.

Begin_transaction Reservation

...

end.

Transaction Structure

- Nested transaction

→ The operations of a transaction may themselves be transactions.

Begin_transaction Reservation

begin

Begin_transaction Airline

 ...

end. {Airline}

Begin_transaction Hotel

 ...

end. {Hotel}

end. {Reservation}

Transactions Provide...

- *Atomic* and *reliable* execution in the presence of failures
- *Correct* execution in the presence of multiple user accesses
- Correct management of *replicas* (if they support it)